

Decoupled Reprojection Consistency for Diagnosing 3D Gaussian Splatting Failures

Jin-Hyeong Park 

Independent Researcher

Abstract

This paper introduces **Decoupled Reprojection Consistency (DRC)**, a training-free diagnostic for 3D Gaussian Splatting (3DGS) that decomposes an ambiguous reprojection error into three interpretable maps: a geometry/visibility mismatch E_{depth} , a base photometric inconsistency E_{base} computed under SH truncation ($L=0$), and a view-dependent residual E_{vd} . By comparing cross-view reprojection errors from zeroth-order (base) and trained-degree (full) SH renders of the same model, DRC separates geometry/visibility-driven inconsistency from appearance-driven instability. Paired with a bivariate fingerprint (quadrant occupancy) and coverage reporting, DRC turns reprojection consistency into actionable per-pixel triage. The method is demonstrated on synthetic and real scenes, separating glossy regions from geometry failures in our examples.

CCS Concepts

• **Computing methodologies** → **Rendering**; Image-based rendering;

1. Introduction

Reprojection consistency—warping one view into another and measuring photometric agreement—is a standard sanity check for radiance-field reconstructions [WBF*17]. However, a single error map is ambiguous: elevated error can stem from geometry/visibility failures (depth errors, occlusions, floaters) *or* from view-dependent appearance that does not warp well (glossy highlights, reflections). In 3DGS [KKLD23], view dependence is encoded via per-Gaussian spherical harmonics (SH), so reflective regions may look plausible per view yet become inconsistent under cross-view warping. Existing training-free consistency tools [STGP26, GRS*24] typically produce a single scalar map per view pair—useful for ranking but ambiguous for debugging.

A simple internal ablation addresses this: rendering the trained model with SH truncated to $L=0$ (*base*) versus the trained SH degree (*full*), and comparing reprojection inconsistencies under both renders to obtain a cause-separating diagnostic axis. Note that $L=0$ is used solely as an internal basis truncation for diagnostics, not as a physical diffuse or albedo decomposition. Concretely, the contributions are: (1) a training-free tri-map decomposition into E_{depth} (geometry/visibility), E_{base} (SH0/base), and E_{vd} (view-dependent residual); (2) coverage-aware reporting with a bivariate fingerprint and quadrant labeling for actionable triage; and (3) demonstrations on synthetic and real scenes from public datasets.

2. Method

Inputs. A trained 3DGS model, calibrated cameras $\{K_i, R_i, t_i\}$, and per-view buffers $\{I_i^{\text{base}}, I_i^{\text{full}}, D_i, A_i\}$. The base render uses SH truncation to the DC term only ($L=0$); the full render uses the model’s trained SH degree. In 3DGS [KKLD23], outgoing radiance for direction $\mathbf{d}$ is $\mathbf{c}(\mathbf{d}) = \sum_{\ell=0}^L \sum_{m=-\ell}^{\ell} \mathbf{c}_{\ell m} Y_{\ell}^m(\mathbf{d})$.

Neighbor selection and warping. For each target view i , neighbors $j \in \mathcal{N}(i)$ are selected via camera-center kNN. Pixels in i are back-projected using D_i to 3D, reprojection into each j , and bilinearly sampled from I_j^{base} and I_j^{full} to obtain warped images $I_j^{\text{base} \rightarrow i}$ and $I_j^{\text{full} \rightarrow i}$.

Masks and coverage. Two validity masks are used per pair (i, j) : (a) M^{geo} : in-bounds $\wedge$ alpha gating—used for E_{depth} ; and (b) M^{photo} : $M^{\text{geo}} \wedge$ depth-consistency—used for E_{base} and E_{full} , where photometric comparison is meaningful only when geometry agrees.

Photometric coverage C_i^{photo} , the fraction of pixels that have at least one valid photometric neighbor, is reported alongside every result.

Error maps. Let $\rho(\cdot, \cdot)$ be the L1 photometric error. For each target view i , inconsistencies across neighbors are robustly aggregated (median) on valid pixels:

- $E_{\text{depth}}^{(i)}(x)$: median relative depth mismatch $|D_j(\bar{x}) - z^{\text{pred}}| / (z^{\text{pred}} + \epsilon)$ across j on M^{geo} .
- $E_{\text{base}}^{(i)}(x)$ and $E_{\text{full}}^{(i)}(x)$: median $\rho(I_i^{\text{base/full}}(x), I_j^{\text{base/full} \rightarrow i}(x))$ across j on M^{photo} .
- $E_{\text{vd}}^{(i)}(x) = \max(0, E_{\text{full}}^{(i)}(x) - E_{\text{base}}^{(i)}(x))$.

Normalization and quadrant labels. E_{base} and E_{full} are normalized by a shared per-view scale $s^{(i)} = q_{90}(\{E_{\text{base}}^{(i)}\} \cup \{E_{\text{full}}^{(i)}\}) + \epsilon$,




Figure 1: DRC pipeline. (A) For target view T and context views $\{C_j\}$, render depth D , base color I_{base} ($L=0$), and full color I_{full} (trained SH degree). (B) Backward-warped context views into T using D_T ; compute E_{depth} , E_{base} , E_{full} , and the residual $E_{vd} = \max(0, E_{full} - E_{base})$. (C) Per-view fingerprint: normalized (E_{base}, E_{full}) scatter with quadrant partitioning.




Figure 2: Qualitative triage. Rows: Mip-NeRF 360 COUNTER (stainless bowl ROI), LLFF ROOM (TV ROI), NeRF-Synthetic TOASTER, LLFF ORCHIDS. Columns: (a) RGB, (b) E_{full} , (c) E_{base} , (d) E_{vd} , (e) quadrant labels. Gray pixels are excluded by masking.

computed on valid pixels, and define $E_{\bullet,n} = E_{\bullet}/s^{(i)}$. With threshold $t=1$, each valid pixel is assigned to one quadrant: *Consistent* (both low), *Geo/Vis-like* (both high), *Appearance-driven* ($E_{base,n} \leq t < E_{full,n}$), or *Debug/Unexpected* ($E_{base,n} > t \geq E_{full,n}$).

3. Experiments

Experiments use Mip-NeRF360 COUNTER, LLFF ROOM and ORCHIDS, and NeRF-Synthetic TOASTER. Fig. 2 shows that the single-map E_{full} conflates photometric mismatch and view-dependent effects: on COUNTER, marble elevates both E_{full} and E_{base} , while the stainless bowl is most prominent in E_{vd} ; on ROOM and TOASTER, glossy surfaces (TV, toaster body) are labeled appearance-driven while edges fall into the Geo/Vis-like regime; on the matte ORCHIDS, appearance-driven pixels are sparse. Table 1 confirms high coverage ($\bar{C}_{photo} \approx 0.85\text{--}0.97$) and that the appearance-driven fraction is consistent with scene glossiness.

Scene (N)	$\bar{C}_{photo}$	Cons.	Geo	App	Dbg
M360/counter (216)	0.947	85.8	5.8	6.1	2.2
LLFF/room (37)	0.967	85.1	5.0	8.1	1.8
Syn/toaster (100)	0.854	85.0	5.0	8.2	1.8
LLFF/orchids (23)	0.852	87.7	7.7	3.1	1.4

Table 1: Coverage and quadrant occupancy on full target views. $\bar{C}_{photo}$: mean photometric coverage. Cons./Geo/App/Dbg: quadrant occupancy (%) on valid pixels with $t=1$.

4. Limitations and Takeaway

DRC is a diagnostic heuristic, not a ground-truth classifier: perfect mirrors and transparent glass can confound depth-based warping, and $L=0$ truncation is not a physical intrinsic decomposition. Low coverage from extreme pose gaps yields sparse maps; therefore, $\bar{C}_{photo}$ is always reported. Despite these caveats, DRC turns an ambiguous single error map into a training-free, coverage-aware tri-map diagnostic that separates geometry/visibility failures from view-dependent instability for practical 3DGS debugging.

References

- [GRS*24] GOLI L., READING C., SELLÁN S., JACOBSON A., TAGLIASACCHI A.: Bayes’ Rays: Uncertainty quantification for neural radiance fields. In *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)* (2024), pp. 20061–20070. doi: 10.1109/CVPR52733.2024.01896. 1
- [KKLD23] KERBL B., KOPANAS G., LEIMKÜHLER T., DRETTAKIS G.: 3D Gaussian splatting for real-time radiance field rendering. *ACM Transactions on Graphics (Proc. SIGGRAPH)* 42, 4 (2023). doi:10.1145/3592433. 1
- [STGP26] SAFADOUST S., TOSI F., GÜNEY F., POGGI M.: WarpRF: Multi-view consistency for training-free uncertainty quantification and applications in radiance fields. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)* (March 2026), pp. 5226–5235. 1
- [WBF*17] WAECHTER M., BELJAN M., FUHRMANN S., MOEHRLE N., KOPF J., GOESELE M.: Virtual rephotography: Novel view prediction error for 3D reconstruction. *ACM Transactions on Graphics* 36, 4 (2017). doi:10.1145/3072959.2999533. 1